

# Deep Learning Project

## Wine Quality Prediction (Multi-class classification)

In [1]:

```
#importing the libraries
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report, r2_score, mean_squared_e
```

In [2]:

```
# call csv file and convert it to dataframe using Pandas function
```

```
df3=pd.read_csv('winequalitynew.csv')
```

```
# .head() will give the first 5 rows of the dataset by default
```

```
df3.head()
```

Out[2]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	

In [3]:

*# .info() will give basic info about the dataset*

df3.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1596 entries, 0 to 1595
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1596 non-null   float64
1   volatile acidity       1596 non-null   float64
2   citric acid            1596 non-null   float64
3   residual sugar         1596 non-null   float64
4   chlorides              1596 non-null   float64
5   free sulfur dioxide    1596 non-null   float64
6   total sulfur dioxide   1596 non-null   float64
7   density                1596 non-null   float64
8   pH                    1596 non-null   float64
9   sulphates              1596 non-null   float64
10  alcohol                1596 non-null   float64
11  quality                1596 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 149.8 KB
```

In [4]:

df3.describe()

Out[4]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total d
<b>count</b>	1596.000000	1596.000000	1596.000000	1596.000000	1596.000000	1596.000000	1596.0
<b>mean</b>	8.314160	0.527954	0.270276	2.535558	0.087120	15.858396	46.3
<b>std</b>	1.732203	0.179176	0.193894	1.405515	0.045251	10.460554	32.8
<b>min</b>	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.0
<b>25%</b>	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.0
<b>50%</b>	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.0
<b>75%</b>	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.0
<b>max</b>	15.600000	1.580000	0.790000	15.500000	0.611000	72.000000	289.0

In [5]:

*# .shape will give number of rows and columns respectively*

df3.shape

Out[5]:

(1596, 12)

## Exploratory Data Analysis

In [6]:

```
# isnull().sum() gives the count of null values present in each column
```

```
df3.isnull().sum()
```

Out[6]:

fixed acidity	0
volatile acidity	0
citric acid	0
residual sugar	0
chlorides	0
free sulfur dioxide	0
total sulfur dioxide	0
density	0
pH	0
sulphates	0
alcohol	0
quality	0

dtype: int64


In [7]:

```
# plot heatmap to find null values present in the dataframe
```

```
sns.heatmap(df3.isnull())
```

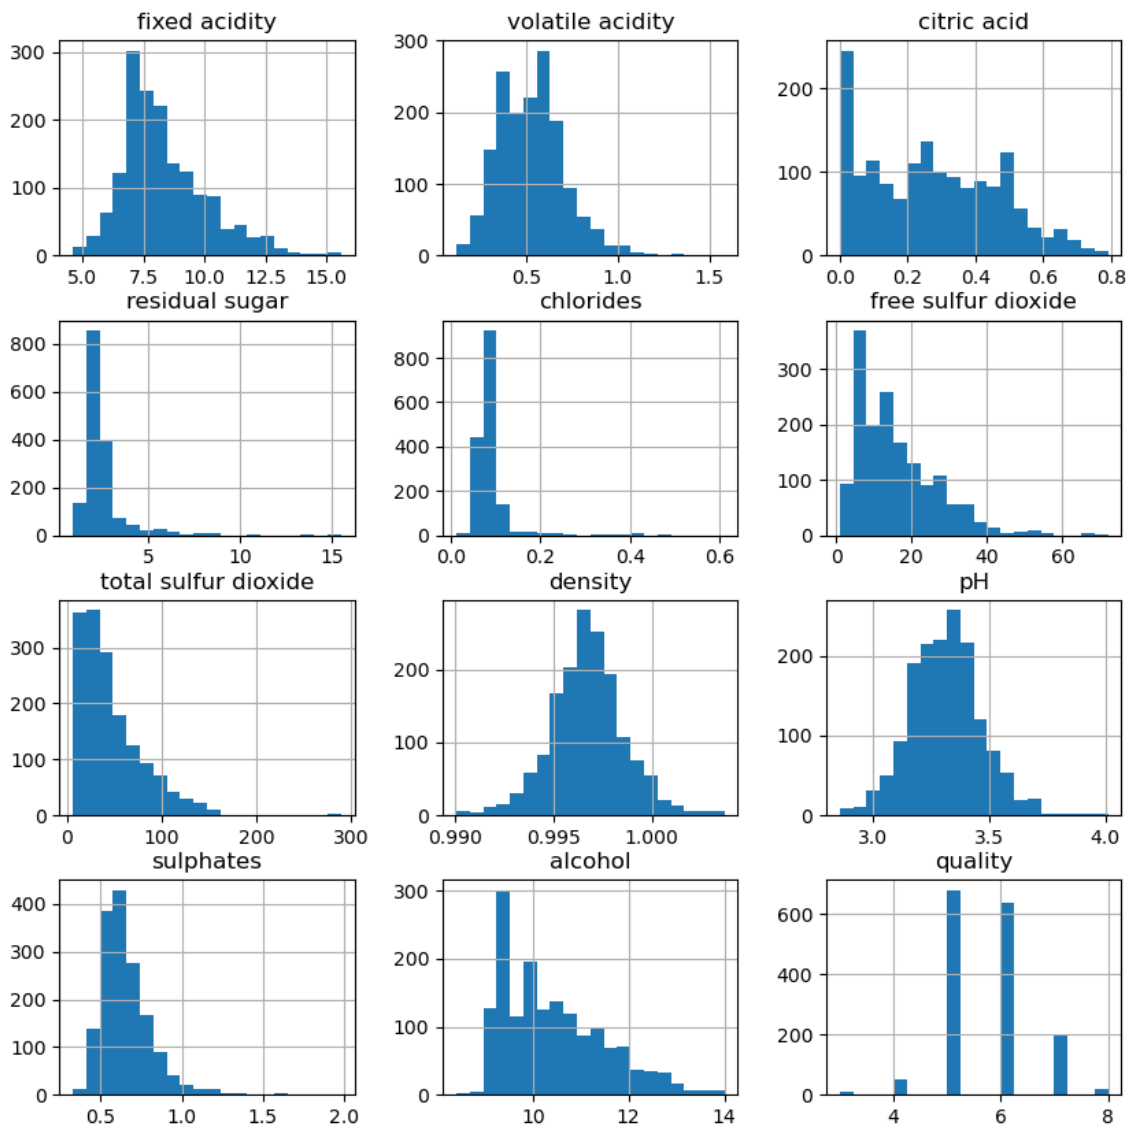
Out[7]:

<AxesSubplot:>



In [8]:

```
# draw the histogram to visualise the distribution of the data with continuous values in  
df3.hist(bins=20, figsize=(10, 10))  
plt.show()
```



In [9]:

```
df3["quality"].unique()
```

Out[9]:


```
array([5, 6, 7, 4, 8, 3], dtype=int64)
```

In [10]:

```
# draw the count plot to visualise the number data for each quality of wine.  
sns.countplot(data=df3,x="quality")
```

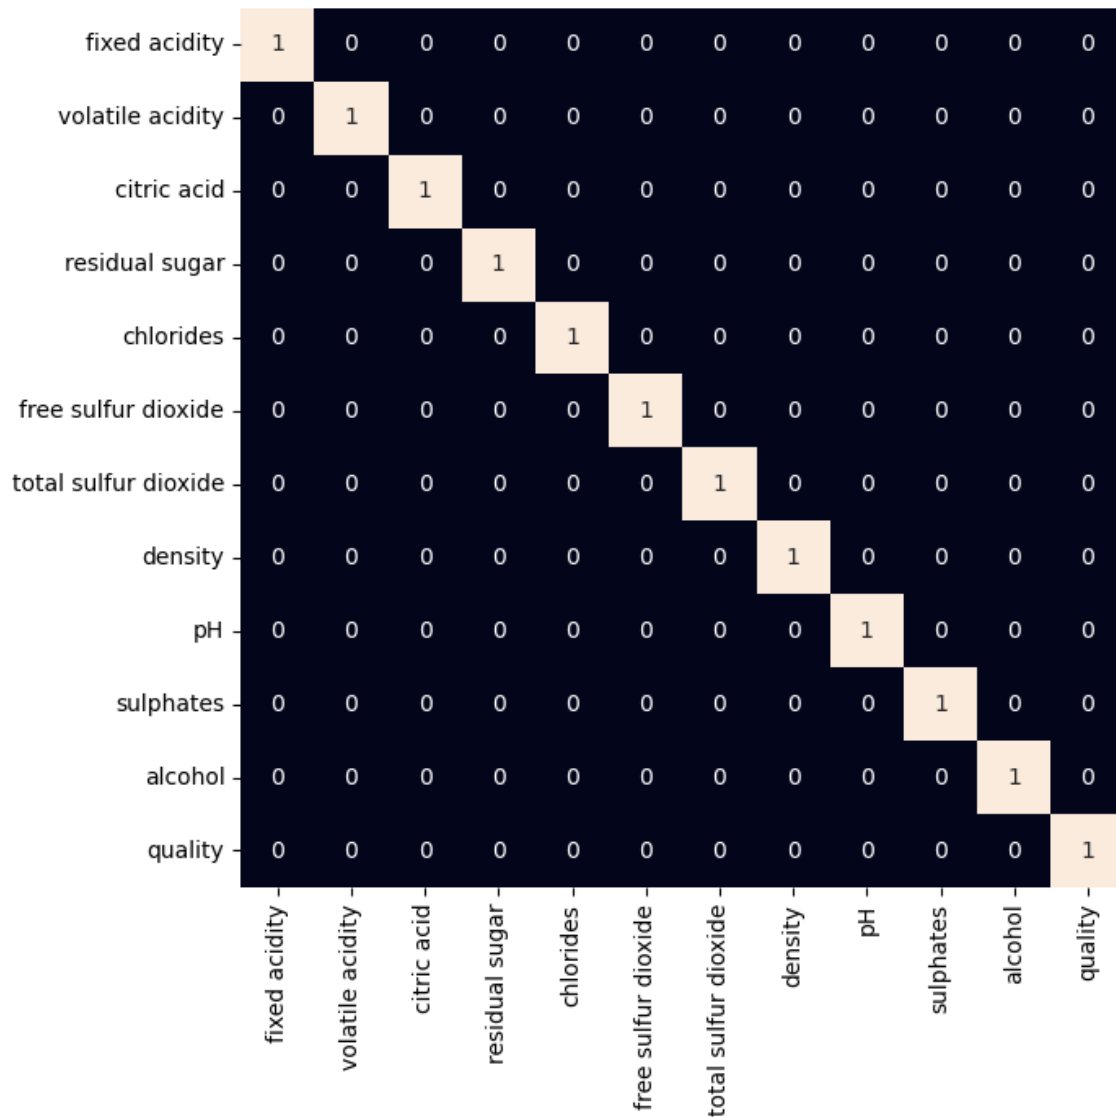
Out[10]:

<AxesSubplot:xlabel='quality', ylabel='count'>



In [11]:

```
plt.figure(figsize=(7, 7))
sns.heatmap(df3.corr() > 0.7, annot=True, cbar=False)
plt.show()
```



In [12]:

```
# split our features and target in x and y respectively
x=df3.iloc[:, :-1]
y=df3.iloc[:, -1]
```

In [13]:

```
x
```

Out[13]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56
...	...	...	...	...	...	...	...	...	...	...
1591	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58
1592	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76
1593	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75
1594	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71
1595	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66

1596 rows × 11 columns

In [14]:

```
y
```

Out[14]:

```
0      5
1      5
2      5
3      6
4      5
..
1591   5
1592   6
1593   6
1594   5
1595   6
Name: quality, Length: 1596, dtype: int64
```

In [15]:

```
from sklearn.preprocessing import LabelEncoder

en=LabelEncoder()
y=en.fit_transform(y)
```



In [16]:

```
y
```

Out[16]:

```
array([2, 2, 2, ..., 3, 2, 3], dtype=int64)
```

In [17]:

```
x.shape
```

Out[17]:

```
(1596, 11)
```

In [18]:

```
sc=StandardScaler()  
x=pd.DataFrame(sc.fit_transform(x),columns=x.columns)  
x.head()
```

Out[18]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	
0	-0.527910	0.960506	-1.394369	-0.452330	-0.245811	-0.464595	-0.377175	0.559577	1.29
1	-0.296918	1.965419	-1.394369	0.045864	0.240520	0.874186	0.628039	0.029638	-0.73
2	-0.296918	1.295477	-1.188007	-0.167648	0.107884	-0.082086	0.232046	0.135626	-0.33
3	1.666516	-1.384292	1.494706	-0.452330	-0.267917	0.109168	0.414812	0.665565	-0.99
4	-0.527910	0.960506	-1.394369	-0.452330	-0.245811	-0.464595	-0.377175	0.559577	1.29

In [19]:

```
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Dense  
from tensorflow.keras.layers import Dropout
```

In [20]:

```
# split the data for training and testing  
xtrain,xtest,ytrain,ytest = train_test_split(x,y,test_size=0.3,random_state=1)
```

In [21]:

```
model=Sequential()  
model.add(Dense(1024,activation="relu",input_dim=11))  
model.add(Dense(512,activation="relu"))  
model.add(Dense(256,activation="relu"))  
model.add(Dense(128,activation="relu"))  
model.add(Dense(64,activation="relu"))  
model.add(Dropout(0.3))  
model.add(Dense(6,activation="softmax"))
```

In [22]:

```
model.compile(optimizer="adam",loss="sparse_categorical_crossentropy",metrics="accuracy")
```

In [23]:

```
his=model.fit(xtrain,ytrain,epochs=50,batch_size=200)
```

```
Epoch 1/50  
6/6 [=====] - 1s 16ms/step - loss: 1.5986 - ac  
curacy: 0.3536  
Epoch 2/50  
6/6 [=====] - 0s 10ms/step - loss: 1.2192 - ac  
curacy: 0.5372  
Epoch 3/50  
6/6 [=====] - 0s 13ms/step - loss: 1.0990 - ac  
curacy: 0.5613  
Epoch 4/50  
6/6 [=====] - 0s 13ms/step - loss: 1.0275 - ac  
curacy: 0.5640  
Epoch 5/50  
6/6 [=====] - 0s 9ms/step - loss: 1.0196 - acc  
uracy: 0.5667  
Epoch 6/50  
6/6 [=====] - 0s 11ms/step - loss: 0.9779 - ac  
curacy: 0.5909  
Epoch 7/50  
6/6 [=====] - 0s 10ms/step - loss: 0.9550 - ac  
curacy: 0.6050
```

In [24]:

```
his.history["loss"]
```

Out[24]:


```
[1.5985742807388306,  
 1.2192264795303345,  
 1.0990239381790161,  
 1.0275427103042603,  
 1.0195618867874146,  
 0.9778688549995422,  
 0.9558177590370178,  
 0.9323015809059143,  
 0.9322930574417114,  
 0.9153008460998535,  
 0.8868661522865295,  
 0.8807466626167297,  
 0.8489855527877808,  
 0.8481696844100952,  
 0.8285951018333435,  
 0.8210949301719666,  
 0.7880850434303284,  
 0.7946338653564453.]
```

In [25]:

```
plt.plot(his.history['loss'])
```

Out[25]:

```
[<matplotlib.lines.Line2D at 0x2b0b8d10430>]
```



In [26]:

```
ypred=model.predict(xtest)
ypred
```

15/15 [=====] - 0s 3ms/step

Out[26]:

```
array([[6.5743268e-05, 2.8209650e-05, 9.9204081e-01, 5.4747961e-03,
        2.3897958e-03, 5.8657469e-07],
       [1.7870902e-05, 4.8415091e-06, 9.4076052e-02, 9.0581959e-01,
        8.1409897e-05, 2.4046389e-07],
       [1.4557431e-03, 1.1085472e-04, 1.8698057e-02, 1.9169832e-03,
        9.7365159e-01, 4.1668084e-03],
       ...,
       [5.7543931e-03, 2.6994050e-03, 8.9250095e-02, 3.0946878e-01,
        1.8659765e-01, 4.0622967e-01],
       [1.2575396e-02, 5.0412379e-02, 9.3272698e-01, 4.0136403e-03,
        2.6332680e-04, 8.1772714e-06],
       [2.8577106e-04, 3.1008727e-05, 9.7925341e-01, 1.1553658e-03,
        1.9254485e-02, 1.9934165e-05]], dtype=float32)
```

In [27]:

```
ypred=np.argmax(ypred,axis=1)
ypred
```

Out[27]:

```
array([2, 3, 4, 3, 3, 4, 4, 4, 2, 3, 2, 2, 4, 4, 3, 2, 3, 3, 2, 2, 2, 0,
        2, 3, 2, 2, 4, 4, 2, 4, 3, 3, 4, 3, 4, 3, 2, 3, 2, 3, 3, 2, 3, 2,
        4, 3, 3, 3, 2, 2, 2, 4, 3, 2, 2, 3, 3, 2, 2, 2, 4, 4, 2, 2, 2, 2,
        4, 3, 2, 1, 3, 2, 3, 3, 2, 2, 3, 4, 3, 3, 2, 2, 3, 2, 2, 2, 2, 2,
        2, 3, 3, 2, 3, 3, 3, 2, 4, 4, 3, 2, 3, 2, 2, 2, 2, 3, 2, 2, 2, 3,
        2, 2, 3, 3, 4, 2, 4, 3, 2, 3, 3, 4, 2, 2, 2, 2, 2, 3, 3, 3, 2, 2,
        3, 2, 2, 2, 3, 2, 2, 3, 3, 4, 3, 2, 3, 3, 2, 2, 3, 2, 2, 3, 2, 3,
        4, 3, 3, 2, 2, 3, 4, 2, 2, 2, 2, 3, 3, 2, 2, 3, 2, 3, 2, 2, 2, 2,
        4, 2, 4, 2, 2, 4, 4, 2, 3, 4, 2, 2, 2, 2, 2, 2, 4, 3, 3, 3, 3, 3,
        3, 2, 2, 3, 2, 3, 3, 2, 2, 3, 2, 3, 2, 3, 2, 3, 3, 3, 2, 3, 2,
        4, 2, 3, 3, 2, 2, 2, 2, 4, 4, 3, 2, 3, 2, 4, 4, 4, 2, 3, 3, 2, 4,
        3, 2, 2, 2, 2, 2, 3, 2, 2, 3, 3, 3, 2, 3, 2, 3, 2, 2, 2, 2, 2,
        2, 2, 3, 3, 3, 2, 2, 3, 2, 2, 3, 3, 2, 2, 2, 2, 2, 2, 3, 3, 3, 1, 4,
        4, 4, 2, 2, 2, 3, 2, 2, 4, 2, 3, 2, 2, 3, 2, 3, 2, 2, 4, 4, 3, 2,
        2, 4, 3, 4, 3, 2, 3, 4, 4, 2, 2, 2, 2, 4, 2, 2, 4, 2, 4, 2, 2, 3,
        1, 3, 4, 2, 2, 3, 3, 2, 2, 4, 2, 2, 4, 2, 2, 2, 2, 2, 2, 2, 2, 3,
        2, 3, 3, 3, 3, 3, 3, 2, 2, 4, 4, 3, 3, 3, 3, 3, 2, 3, 2, 2, 2, 4,
        2, 3, 3, 3, 2, 2, 4, 0, 3, 2, 3, 3, 2, 4, 4, 2, 2, 1, 2, 2, 3, 2,
        3, 3, 2, 3, 2, 2, 4, 4, 4, 4, 2, 3, 2, 2, 3, 3, 2, 3, 2, 2, 2, 3,
        2, 2, 4, 3, 2, 3, 3, 3, 4, 2, 3, 4, 2, 3, 2, 4, 3, 2, 2, 2, 3, 2,
        3, 3, 3, 3, 3, 3, 3, 3, 4, 2, 4, 3, 3, 4, 5, 2, 2], dtype=int64)
```

In [28]:

```
print(classification_report(ytest,ypred))
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	2
1	0.00	0.00	0.00	18
2	0.72	0.78	0.75	212
3	0.66	0.60	0.63	189
4	0.49	0.65	0.56	54
5	0.00	0.00	0.00	4
accuracy			0.65	479
macro avg	0.31	0.34	0.32	479
weighted avg	0.63	0.65	0.64	479

In [29]:

```
print(accuracy_score(ytest,ypred))
```

0.6534446764091858